

Part-2:

Copy the attached 'access.log' file into HDFS under /logs directory.

Using the access.log file stored in HDFS, implement MapReduce in Hadoop to find the number of times each IP accessed the website.

Run the job without a Combiner, and see how long it will take (you may use your clock or smart watch to find the running time)

Then, add the Combiner and see if it will run faster. Since Counting is both commutative and associative operation, you could use the same Reducer as a Combiner.

```
[cloudera@quickstart Desktop]$ hadoop fs -mkdir /logs
[cloudera@quickstart Desktop]$ hadoop fs -copyFromLocal /home/cloudera/Desktop/hw4/access.log /logs/
[cloudera@quickstart Desktop]$
```

Hadoop

Overview

Datanodes

Snapshot

Startup Progress

Utilities

Browse Directory

/logs

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	cloudera	supergroup	3.34 MB	Sun Oct 30 22:02:18 -0700 2022	1	128 MB	access.log

Hadoop, 2017.

Mapper Class:

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import java.io.IOException;

public class IPAddrMapper extends Mapper<Object, Text, Text, IntWritable>{
    private final static IntWritable one = new IntWritable(1);
    Text ipAddressText = new Text();
    //called once for each key/value pair in the input split
    protected void map(Object key, Text value, Mapper<Object, Text, Text,
```

```

IntWritable>.Context context) throws IOException, InterruptedException {

    String line = value.toString();

    String[] tokens = line.split(" ");

    String ipAddress = tokens[0].trim();

    System.out.println("Mapping ipAddress: " + ipAddress);

    ipAddressText.set(ipAddress);

    context.write(ipAddressText, one); //emit(key,value)

}
}

```

Reducer Class:

```

public class IPAddrReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    protected void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable,
        Text, IntWritable>.Context context) throws IOException, InterruptedException {

        IntWritable sumObj = new IntWritable(0);

        int sum = 0;

        for(IntWritable val:values) {

            sum += val.get();

        }

        sumObj.set(sum);

        System.out.println("Reducing key: "+ key + " Final count: " + sum);

        context.write(key, sumObj); //returning Hadoop Datatype

    }

}

```

Driver Class:

```

public class IPAddrDriver {

    public static void main(String[] args) throws IOException, InterruptedException,
        ClassNotFoundException {

        Configuration configuration = new Configuration();
    }
}

```

```

Job job = Job.getInstance(configuration, "IP Address COUNTER");

job.setJarByClass(IPAddrDriver.class);

job.setMapperClass(IPAddrMapper.class);

job.setReducerClass(IPAddrReducer.class);

//Reducer Output

job.setOutputKeyClass(Text.class);

job.setOutputValueClass(IntWritable.class);

FileInputFormat.addInputPath(job, new Path(args[0]));

FileOutputFormat.setOutputPath(job, new Path(args[1]));

System.exit(job.waitForCompletion(true) ? 0 : 1);

}

}

```

Run Command:

hadoop jar /home/cloudera/Desktop/hw4/checkip.jar IPAddrDriver /logs/access.log /hw5output

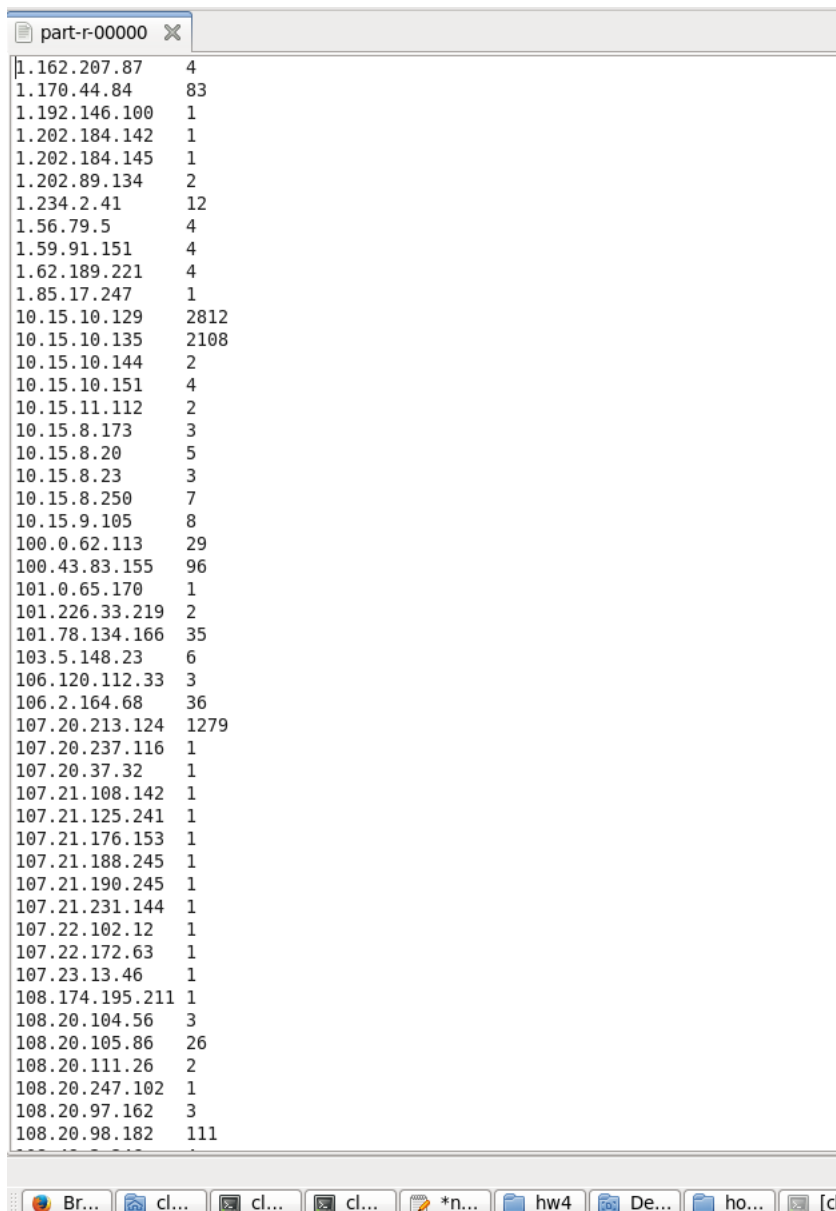
```

[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Desktop/hw4/checkip.jar IPAddrDriver /logs/access.log /hw5output
22/10/30 22:24:05 INFO client.RMProxy: Connecting to ResourceManager at /b.d.b.d:8032
22/10/30 22:24:06 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/10/30 22:24:06 INFO InputFileInputFormat: Total input paths to process : 1
22/10/30 22:24:06 INFO mapreduce.JobSubmitter: number of splits:1
22/10/30 22:24:06 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1666102279738_0001
22/10/30 22:24:07 INFO impl.YarnClientImpl: Submitted application application_1666102279738_0001
22/10/30 22:24:07 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8080/proxy/application_1666102279738_0001/
22/10/30 22:24:07 INFO mapreduce.Job: Running job: job_1666102279738_0001
22/10/30 22:24:15 INFO mapreduce.Job: Job job_1666102279738_0001 running in uber mode : false
22/10/30 22:24:15 INFO mapreduce.Job:  map 0% reduce 0%
22/10/30 22:24:28 INFO mapreduce.Job:  map 100% reduce 0%
22/10/30 22:24:36 INFO mapreduce.Job:  map 100% reduce 100%
22/10/30 22:24:37 INFO mapreduce.Job: Job job_1666102279738_0001 completed successfully
22/10/30 22:24:38 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=709513
    FILE: Number of bytes written=1608653
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=3497779
    HDFS: Number of bytes written=31931
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2

```

Run Time: 25 seconds

Output:



1.162.207.87	4
1.170.44.84	83
1.192.146.100	1
1.202.184.142	1
1.202.184.145	1
1.202.89.134	2
1.234.2.41	12
1.56.79.5	4
1.59.91.151	4
1.62.189.221	4
1.85.17.247	1
10.15.10.129	2812
10.15.10.135	2108
10.15.10.144	2
10.15.10.151	4
10.15.11.112	2
10.15.8.173	3
10.15.8.20	5
10.15.8.23	3
10.15.8.250	7
10.15.9.105	8
100.0.62.113	29
100.43.83.155	96
101.0.65.170	1
101.226.33.219	2
101.78.134.166	35
103.5.148.23	6
106.120.112.33	3
106.2.164.68	36
107.20.213.124	1279
107.20.237.116	1
107.20.37.32	1
107.21.108.142	1
107.21.125.241	1
107.21.176.153	1
107.21.188.245	1
107.21.190.245	1
107.21.231.144	1
107.22.102.12	1
107.22.172.63	1
107.23.13.46	1
108.174.195.211	1
108.20.104.56	3
108.20.105.86	26
108.20.111.26	2
108.20.247.102	1
108.20.97.162	3
108.20.98.182	111

With Combiner:

Driver Class

```
public class IPAddrDriver {  
  
    public static void main(String[] args) throws IOException, InterruptedException,  
    ClassNotFoundException {  
  
        Configuration configuration = new Configuration();  
  
        Job job = Job.getInstance(configuration, "IP Address COUNTER");
```

```

        job.setJarByClass(IPAddrDriver.class);

        job.setMapperClass(IPAddrMapper.class);

        job.setCombinerClass(IPAddrReducer.class);

        job.setReducerClass(IPAddrReducer.class);

        //Reducer Output

        job.setOutputKeyClass(Text.class);

        job.setOutputValueClass(IntWritable.class);


        FileInputFormat.addInputPath(job, new Path(args[0]));

        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        System.exit(job.waitForCompletion(true) ? 0 : 1);

    }

}

```

Run Command:

`hadoop jar /home/cloudera/Desktop/hw4/combiner.jar IPAddrDriver /logs/access.log /output1`

```

[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Desktop/hw4/combiner.jar IPAddrDriver /logs/access.log /output1
22/10/30 22:56:52 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/10/30 22:56:52 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/10/30 22:56:53 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:952)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:690)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:879)
22/10/30 22:56:53 INFO input.FileInputFormat: Total input paths to process : 1
22/10/30 22:56:53 INFO mapreduce.JobSubmitter: number of splits:1
22/10/30 22:56:53 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1666102279738_0003
22/10/30 22:56:53 INFO impl.YarnClientImpl: Submitted application application_1666102279738_0003
22/10/30 22:56:53 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1666102279738_0003/
22/10/30 22:56:53 INFO mapreduce.Job: Running job: job_1666102279738_0003
22/10/30 22:56:59 INFO mapreduce.Job: Job job_1666102279738_0003 running in uber mode : false
22/10/30 22:56:59 INFO mapreduce.Job:  map 0% reduce 0%
22/10/30 22:57:06 INFO mapreduce.Job:  map 100% reduce 0%
22/10/30 22:57:13 INFO mapreduce.Job:  map 100% reduce 100%
22/10/30 22:57:13 INFO mapreduce.Job: Job job_1666102279738_0003 completed successfully
22/10/30 22:57:14 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=39401
  FILE: Number of bytes written=328725
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=3497779
  HDFS: Number of bytes written=31931
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2

```

Run Time: 22 seconds

Output:

part-r-00000	part-r-00000(1)
1.162.207.87	4
1.170.44.84	83
1.192.146.100	1
1.202.184.142	1
1.202.184.145	1
1.202.89.134	2
1.234.2.41	12
1.56.79.5	4
1.59.91.151	4
1.62.189.221	4
1.85.17.247	1
10.15.10.129	2812
10.15.10.135	2108
10.15.10.144	2
10.15.10.151	4
10.15.11.112	2
10.15.8.173	3
10.15.8.20	5
10.15.8.23	3
10.15.8.250	7
10.15.9.105	8
100.0.62.113	29
100.43.83.155	96
101.0.65.170	1
101.226.33.219	2
101.78.134.166	35
103.5.148.23	6
106.120.112.33	3
106.2.164.68	36
107.20.213.124	1279
107.20.237.116	1
107.20.37.32	1
107.21.108.142	1
107.21.125.241	1
107.21.176.153	1
107.21.188.245	1
107.21.190.245	1
107.21.231.144	1
107.22.102.12	1
107.22.172.63	1
107.23.13.46	1
108.174.195.211	1
108.20.104.56	3
108.20.105.86	26
108.20.111.26	2
108.20.247.102	1
108.20.97.162	3
108.20.98.182	111

Part-3: Programming Assignment

Download and Copy all the files (<http://newton.neu.edu/nyse/> (Links to an external site.))

```
public class copyNYSE {  
  
    public static void main(String[] args) throws Exception {  
  
        Configuration config = new Configuration();  
  
        FileSystem hdfs = FileSystem.get(config);  
  
        FileSystem local = FileSystem.getLocal(config);
```

```

for(char i = 'A' ; i <= 'Z'; i++) {

    String csvPathPrefix = "NYSE_daily_prices_" + i;

    String fullPath =

        "/home/cloudera/Desktop/shared_nyse/" +

        csvPathPrefix + ".csv";

    System.out.println(fullPath);

    hdfs.copyFromLocalFile(new Path(fullPath), new Path("/hw4/NYSE/"));

}

}

```

Run Command: `hadoop jar /home/cloudera/Desktop/hw4/copyNYSE.jar NYSE.copyNYSE /shared_nyse /output2`

```

at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Desktop/hw4/copyNYSE.jar NYSE.copyNYSE /shared_nyse /output2
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_A.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_B.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_C.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_D.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_E.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_F.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_G.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_H.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_I.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_J.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_K.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_L.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_M.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_N.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_O.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_P.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_Q.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_R.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_S.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_T.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_U.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_V.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_W.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_X.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_Y.csv
/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_Z.csv
[cloudera@quickstart ~]$

```

Browse Files:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	cloudera	supergroup	39.09 MB	Mon Oct 31 00:09:18 -0700 2022	1	128 MB	NYSE_daily_prices_A.csv
-rw-r--r--	cloudera	supergroup	30.55 MB	Mon Oct 31 00:09:19 -0700 2022	1	128 MB	NYSE_daily_prices_B.csv
-rw-r--r--	cloudera	supergroup	43.67 MB	Mon Oct 31 00:09:20 -0700 2022	1	128 MB	NYSE_daily_prices_C.csv
-rw-r--r--	cloudera	supergroup	18.34 MB	Mon Oct 31 00:09:20 -0700 2022	1	128 MB	NYSE_daily_prices_D.csv
-rw-r--r--	cloudera	supergroup	21.08 MB	Mon Oct 31 00:09:20 -0700 2022	1	128 MB	NYSE_daily_prices_E.csv
-rw-r--r--	cloudera	supergroup	16.58 MB	Mon Oct 31 00:09:20 -0700 2022	1	128 MB	NYSE_daily_prices_F.csv
-rw-r--r--	cloudera	supergroup	21.56 MB	Mon Oct 31 00:09:21 -0700 2022	1	128 MB	NYSE_daily_prices_G.csv
-rw-r--r--	cloudera	supergroup	22.06 MB	Mon Oct 31 00:09:21 -0700 2022	1	128 MB	NYSE_daily_prices_H.csv
-rw-r--r--	cloudera	supergroup	19.72 MB	Mon Oct 31 00:09:22 -0700 2022	1	128 MB	NYSE_daily_prices_I.csv
-rw-r--r--	cloudera	supergroup	9.1 MB	Mon Oct 31 00:09:22 -0700 2022	1	128 MB	NYSE_daily_prices_J.csv
-rw-r--r--	cloudera	supergroup	14.1 MB	Mon Oct 31 00:09:22 -0700 2022	1	128 MB	NYSE_daily_prices_K.csv
-rw-r--r--	cloudera	supergroup	12.36 MB	Mon Oct 31 00:09:22 -0700 2022	1	128 MB	NYSE_daily_prices_L.csv
-rw-r--r--	cloudera	supergroup	36.36 MB	Mon Oct 31 00:09:23 -0700 2022	1	128 MB	NYSE_daily_prices_M.csv
-rw-r--r--	cloudera	supergroup	30.03 MB	Mon Oct 31 00:09:23 -0700 2022	1	128 MB	NYSE_daily_prices_N.csv
-rw-r--r--	cloudera	supergroup	8.46 MB	Mon Oct 31 00:09:24 -0700 2022	1	128 MB	NYSE_daily_prices_O.csv
-rw-r--r--	cloudera	supergroup	30.46 MB	Mon Oct 31 00:09:24 -0700 2022	1	128 MB	NYSE_daily_prices_P.csv
-rw-r--r--	cloudera	supergroup	186.51 KB	Mon Oct 31 00:09:24 -0700 2022	1	128 MB	NYSE_daily_prices_Q.csv
-rw-r--r--	cloudera	supergroup	16.03 MB	Mon Oct 31 00:09:24 -0700 2022	1	128 MB	NYSE_daily_prices_R.csv
-rw-r--r--	cloudera	supergroup	30.38 MB	Mon Oct 31 00:09:25 -0700 2022	1	128 MB	NYSE_daily_prices_S.csv
-rw-r--r--	cloudera	supergroup	27.42 MB	Mon Oct 31 00:09:25 -0700 2022	1	128 MB	NYSE_daily_prices_T.csv

3.1: Copy NYSE dataset (DailyPrices_A to DailyPrices_Z) to a folder in HDFS. No merging.

Write a MapReduce to find the Max price of stock_price_high for each stock. Capture the running time programmatically (or manually using a wristwatch or smartphone).

Mapper Class:

```
public class MaxStockPriceMap extends Mapper<Object, Text, Text, DoubleWritable> {

    protected void map(Object key, Text value, Mapper<Object, Text, Text,
        DoubleWritable>.Context context) throws IOException, InterruptedException {

        String[] tokens = value.toString().split(",");

        Text stockSymbol = new Text();

        stockSymbol.set(tokens[1]);

        DoubleWritable price = new DoubleWritable();

        Double stock_price_high_value = 0.0;

        try {

            stock_price_high_value = Double.parseDouble(tokens[4]);

            price.set(stock_price_high_value);

        }
    }
}
```



```

        catch (NumberFormatException nf) {
            System.out.println("Continue without interruption, its the header line");
        }
        context.write(stockSymbol, price);
    }
}
}

```

Reducer Class:

```

public class MaxStockPriceReducer extends Reducer<Text, DoubleWritable, Text, DoubleWritable> {
    public void reduce(Text key, Iterable<DoubleWritable> values, Reducer.Context context) throws
    IOException, InterruptedException {
        double final_stock_high = Double.MIN_VALUE;
        for (DoubleWritable val: values) {
            final_stock_high = Math.max(val.get(), final_stock_high);
        }
        DoubleWritable final_stock_price_red = new DoubleWritable(final_stock_high);
        context.write(key, final_stock_price_red);
    }
}

```

Driver Class:

```

public class MaxStockPriceDriver {
    public static void main(String[] args) throws Exception {
        Configuration configuration = new Configuration();
        Job job = Job.getInstance(configuration, "Calculating max stock price high");
        job.setJarByClass(MaxStockPriceDriver.class);
        job.setMapperClass(MaxStockPriceMap.class);
        job.setCombinerClass(MaxStockPriceReducer.class);
        job.setReducerClass(MaxStockPriceReducer.class);
        //output of the mapper
    }
}

```

```

job.setOutputKeyClass(Text.class);
job.setOutputValueClass(DoubleWritable.class);
StringBuilder fullPath = new StringBuilder("");
for(char i = 'A' ; i <= 'Z'; i++) {
    String csvPathPrefix = "NYSE_daily_prices_" + i + ".csv";
    if (i == 'Z') {
        fullPath.append(args[0]).append(csvPathPrefix);
    }
    else {
        fullPath.append(args[0]).append(csvPathPrefix).append(",");
    }
}

String finalcommaSeperatedPath = fullPath.toString();
System.out.println(finalcommaSeperatedPath);
FileInputFormat.addInputPaths(job, finalcommaSeperatedPath);
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Run Command:

```
hadoop jar /home/cloudera/Desktop/hw4/NYSEmerge.jar NYSEmerge.MaxStockPriceDriver /nyse/
/output4
```

Run Time: 2 minutes 8 seconds

```
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Desktop/hw4/NYSEmerge.jar NYSEmerge.MaxStockPriceDriver /nyse/ /output4
/nyse/NYSE_daily_prices_A.csv,/nyse/NYSE_daily_prices_B.csv,/nyse/NYSE_daily_prices_C.csv,/nyse/NYSE_daily_prices_D.csv,/nyse/NYSE_daily_prices_E.csv,/nyse/NYSE_daily_prices_F.csv,/nyse/NYSE_daily_prices_G.csv,/nyse/NYSE_daily_prices_H.csv,/nyse/NYSE_daily_prices_I.csv,/nyse/NYSE_daily_prices_J.csv,/nyse/NYSE_daily_prices_K.csv,/nyse/NYSE_daily_prices_L.csv,/nyse/NYSE_daily_prices_M.csv,/nyse/NYSE_daily_prices_N.csv,/nyse/NYSE_daily_prices_O.csv,/nyse/NYSE_daily_prices_P.csv,/nyse/NYSE_daily_prices_Q.csv,/nyse/NYSE_daily_prices_R.csv,/nyse/NYSE_daily_prices_S.csv,/nyse/NYSE_daily_prices_T.csv,/nyse/NYSE_daily_prices_U.csv,/nyse/NYSE_daily_prices_V.csv,/nyse/NYSE_daily_prices_W.csv,/nyse/NYSE_daily_prices_X.csv,/nyse/NYSE_daily_prices_Y.csv,/nyse/NYSE_daily_prices_Z.csv
22/10/31 00:39:20 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/10/31 00:39:20 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/10/31 00:39:21 INFO input.FileInputFormat: Total input paths to process : 26
22/10/31 00:39:21 INFO mapreduce.JobSubmitter: number of splits:26
22/10/31 00:39:21 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1666102279738_0005
22/10/31 00:39:21 INFO Impl.YarnClientImpl: Submitted application application_1666102279738_0005
22/10/31 00:39:21 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1666102279738_0005/
22/10/31 00:39:21 INFO mapreduce.Job: Running job: job_1666102279738_0005
22/10/31 00:39:28 INFO mapreduce.Job: Job job_1666102279738_0005 running in uber mode : false
22/10/31 00:39:28 INFO mapreduce.Job: map 0% reduce 0%
22/10/31 00:39:57 INFO mapreduce.Job: map 6% reduce 0%
22/10/31 00:39:58 INFO mapreduce.Job: map 12% reduce 0%
22/10/31 00:40:12 INFO mapreduce.Job: map 19% reduce 0%
22/10/31 00:40:13 INFO mapreduce.Job: map 23% reduce 0%
22/10/31 00:40:26 INFO mapreduce.Job: map 35% reduce 0%
22/10/31 00:40:41 INFO mapreduce.Job: map 42% reduce 0%
22/10/31 00:40:46 INFO mapreduce.Job: map 42% reduce 14%
22/10/31 00:40:52 INFO mapreduce.Job: map 54% reduce 14%
22/10/31 00:40:58 INFO mapreduce.Job: map 58% reduce 18%
22/10/31 00:41:00 INFO mapreduce.Job: map 62% reduce 18%
22/10/31 00:41:04 INFO mapreduce.Job: map 62% reduce 21%
22/10/31 00:41:16 INFO mapreduce.Job: map 65% reduce 21%
22/10/31 00:41:17 INFO mapreduce.Job: map 73% reduce 21%
22/10/31 00:41:19 INFO mapreduce.Job: map 77% reduce 21%
22/10/31 00:41:20 INFO mapreduce.Job: map 81% reduce 21%
22/10/31 00:41:22 INFO mapreduce.Job: map 81% reduce 27%
22/10/31 00:41:38 INFO mapreduce.Job: map 92% reduce 27%
22/10/31 00:41:39 INFO mapreduce.Job: map 100% reduce 27%
22/10/31 00:41:40 INFO mapreduce.Job: map 100% reduce 100%
22/10/31 00:41:42 INFO mapreduce.Job: Job job_1666102279738_0005 completed successfully
22/10/31 00:41:42 INFO mapreduce.Job: Counters: 50
  File System Counters
    FILE: Number of bytes read=604
    FILE: Number of bytes written=3419681
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=511088077
    HDFS: Number of bytes written=442
    HDFS: Number of read operations=81
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Killed map tasks=1
    Launched map tasks=26
    Launched reduce tasks=1
```

Output:

AA	94.62
AAI	57.88
AAN	35.21
AAP	83.65
AAR	25.25
AAV	24.78
AB	94.94
ABA	27.94
ABB	33.39
ABC	84.35
ABD	28.58
ABG	38.06
ABK	96.1
ABM	41.63
ABR	34.45
ABT	93.37
ABV	187.5
ABVT	188.0
ABX	54.74
ACC	37.0
ACE	184.0
ACF	64.9
ACG	12.63
ACH	111.6
ACT	112.89
ACL	178.56
ACM	38.25
ACN	44.03
ACO	42.7
ACS	189.55
ACV	65.32
ADC	37.7
ADT	185.5
ADM	48.95
ADP	84.51
ADS	86.79
ADX	48.56
ADY	44.0
AEA	23.94
AEB	26.5
AEC	17.6
AED	26.12
AEE	56.77
AFF	27.0
AEG	148.32
AEH	26.64
AEL	14.6
AEI	83.45
AEO	88.13
AEP	53.31
AER	32.82
AES	92.5
AET	154.67
AEV	26.78
AF	63.09
AFB	17.03
AFC	25.15
AFE	26.7
AFF	25.15
AFG	54.65
AFL	74.94
AFN	11.99
AGC	28.2
AGCO	71.95
AGD	25.5
AGL	44.67
AGM	88.0
AGN	125.0
AGO	31.99
AGP	88.89
AGU	113.88
AHC	16.35
AHD	47.12

3.2: Redo 3.1. Merge the NYSE files in a single file (that you may have done in your prev. HW) on HDFS. Now, repeat 3.1 on the single merged-file. Capture the running time. Did MapReduce on a single file run faster than running MapReduce on a bunch of files?

Part-A:

Code for merging NYSE files:

```
public class MergeNYSE {

    public static void main(String args[]) throws Exception {

        Configuration config = new Configuration();

        FileSystem hdfs = FileSystem.get(config);

        FileSystem local = FileSystem.getLocal(config);
```

```

Path op = new Path(args[1]);
Path ip = new Path(args[0]);
try {
    FileStatus[] ipFile = local.listStatus(ip);
    FSDataOutputStream outputStream = hdfs.create(op);
    System.out.println("Total files " + ipFile.length);
    for (int i = 0; i < ipFile.length; i++) {
        System.out.println("files " + ipFile[i].getPath());
        FSDataInputStream input = local.open(ipFile[i].getPath());
        byte buffer[] = new byte[256];
        int bytesRead = 0;
        while ((bytesRead = input.read(buffer)) > 0) {
            outputStream.write(buffer, 0, bytesRead);
        }
        input.close();
    }
    outputStream.close();
} catch (IOException e) {
    e.printStackTrace();
}
}
}

```

Run Command:

```

hadoop jar /home/cloudera/Desktop/hw4/q3p2.jar mergeAll.MergeNYSE
/home/cloudera/Desktop/shared_nyse /nyseMerged

```

```

at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Desktop/hw4/q3p2.jar mergeAll.MergeNYSE /home/cloudera/Desktop/shared_nyse /nyseMerged
Total files 26
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_A.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_B.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_C.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_D.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_E.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_F.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_G.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_H.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_I.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_J.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_K.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_L.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_M.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_N.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_O.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_P.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_Q.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_R.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_S.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_T.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_U.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_V.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_W.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_X.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_Y.csv
files file:/home/cloudera/Desktop/shared_nyse/NYSE_daily_prices_Z.csv
[cloudera@quickstart ~]$ █

```

Merged Files:

Browse Directory

/								Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
drwxrwxrwx	hdfs	supergroup	0 B	Wed Jul 19 05:34:46 -0700 2017	0	0 B	benchmarks	
drwxr-xr-x	cloudera	supergroup	0 B	Tue Sep 20 12:58:19 -0700 2022	0	0 B	data	
drwxr-xr-x	hbase	supergroup	0 B	Tue Oct 18 09:59:14 -0700 2022	0	0 B	hbase	
drwxr-xr-x	cloudera	supergroup	0 B	Mon Oct 31 00:04:43 -0700 2022	0	0 B	hw4	
drwxr-xr-x	cloudera	supergroup	0 B	Sun Oct 30 22:24:35 -0700 2022	0	0 B	hw5output	
drwxr-xr-x	cloudera	supergroup	0 B	Sun Oct 30 22:31:25 -0700 2022	0	0 B	hw6output	
drwxr-xr-x	cloudera	supergroup	0 B	Sun Oct 30 22:02:18 -0700 2022	0	0 B	logs	
drwxr-xr-x	cloudera	supergroup	0 B	Mon Oct 31 00:09:26 -0700 2022	0	0 B	nyse	
-rw-r--r--	cloudera	supergroup	487.41 MB	Mon Oct 31 06:23:02 -0700 2022	1	128 MB	nyseMerged	
drwxr-xr-x	cloudera	supergroup	0 B	Sun Oct 30 22:57:12 -0700 2022	0	0 B	output1	
drwxr-xr-x	cloudera	supergroup	0 B	Mon Oct 31 00:41:39 -0700 2022	0	0 B	output4	
drwxr-xr-x	cloudera	supergroup	0 B	Mon Oct 31 06:02:42 -0700 2022	0	0 B	output5	

Output:

```
part-r-00000  x  part-r-00000(1)  x  nyseMerged  x
exchange,stock_symbol,date,stock_price_open,stock_price_high,stock_price_low,stock_price_close,stock_volume,stock_price_adj_close
NYSE,AEA,2010-02-08,4.42,4.42,4.21,4.24,205500,4.24
NYSE,AEA,2010-02-05,4.42,4.54,4.22,4.41,194300,4.41
NYSE,AEA,2010-02-04,4.55,4.69,4.39,4.42,233800,4.42
NYSE,AEA,2010-02-03,4.65,4.69,4.50,4.55,182100,4.55
NYSE,AEA,2010-02-02,4.74,5.00,4.62,4.66,222700,4.66
NYSE,AEA,2010-02-01,4.84,4.92,4.68,4.75,194800,4.75
NYSE,AEA,2010-01-29,4.97,5.05,4.76,4.83,222900,4.83
NYSE,AEA,2010-01-28,5.12,5.22,4.81,4.98,283100,4.98
NYSE,AEA,2010-01-27,4.82,5.16,4.79,5.09,243500,5.09
NYSE,AEA,2010-01-26,5.18,5.18,4.81,4.84,554800,4.84
NYSE,AEA,2010-01-25,5.42,5.48,5.20,5.22,257300,5.22
NYSE,AEA,2010-01-22,5.52,5.59,5.31,5.37,260800,5.37
NYSE,AEA,2010-01-21,5.67,5.74,5.37,5.51,264300,5.51
NYSE,AEA,2010-01-20,5.65,5.70,5.53,5.66,244600,5.66
NYSE,AEA,2010-01-19,5.54,5.70,5.54,5.69,368000,5.69
NYSE,AEA,2010-01-15,5.48,5.55,5.33,5.54,435500,5.54
NYSE,AEA,2010-01-14,5.41,5.50,5.39,5.41,272200,5.41
NYSE,AEA,2010-01-13,5.50,5.50,5.41,5.45,176400,5.45
NYSE,AEA,2010-01-12,5.47,5.51,5.41,5.46,233100,5.46
NYSE,AEA,2010-01-11,5.64,5.64,5.49,5.55,178900,5.55
NYSE,AEA,2010-01-08,5.61,5.68,5.52,5.59,144200,5.59
NYSE,AEA,2010-01-07,5.47,5.65,5.40,5.62,228900,5.62
NYSE,AEA,2010-01-06,5.56,5.70,5.44,5.49,208900,5.49
NYSE,AEA,2010-01-05,5.55,5.62,5.51,5.55,267000,5.55
NYSE,AEA,2010-01-04,5.65,5.66,5.49,5.55,335500,5.55
NYSE,AEA,2009-12-31,5.57,5.71,5.54,5.56,418600,5.56
NYSE,AEA,2009-12-30,5.65,5.67,5.50,5.57,226400,5.57
NYSE,AEA,2009-12-29,5.67,5.74,5.66,5.67,115100,5.67
NYSE,AEA,2009-12-28,5.81,5.86,5.63,5.67,326600,5.67
NYSE,AEA,2009-12-24,5.92,5.94,5.81,5.84,111900,5.84
NYSE,AEA,2009-12-23,5.91,5.99,5.84,5.87,212000,5.87
NYSE,AEA,2009-12-22,5.99,6.10,5.84,5.92,307500,5.92
NYSE,AEA,2009-12-21,6.00,6.20,5.90,5.99,257700,5.99
```

Part-B:

Mapper Class:

```
public class MaxStockPriceMap extends Mapper<Object, Text, Text, DoubleWritable> {

    protected void map(Object key, Text value, Mapper<Object, Text, Text,
        DoubleWritable>.Context context) throws IOException, InterruptedException {

        String[] tokens = value.toString().split(",");

        Text stockSymbol = new Text();

        stockSymbol.set(tokens[1]);

        DoubleWritable price = new DoubleWritable();

        Double stock_price_high_value = 0.0;

        try {

            stock_price_high_value = Double.parseDouble(tokens[4]);

            price.set(stock_price_high_value);

        }

        catch (NumberFormatException nf) {

            System.out.println("Continue without interruption, its the header line");
```

```

    }

    context.write(stockSymbol, price);
}

}

```

Reducer Class:

```

public class MaxStockPriceReduce extends Reducer<Text, DoubleWritable, Text, DoubleWritable> {

    public void reduce(Text key, Iterable<DoubleWritable> values, Context context)
        throws IOException, InterruptedException {

        double final_stock_high = Double.MIN_VALUE;

        for (DoubleWritable val: values) {

            final_stock_high = Math.max(val.get(), final_stock_high);

        }

        DoubleWritable final_stock_price_red = new DoubleWritable(final_stock_high);

        context.write(key, final_stock_price_red);

    }

}

```

Driver Class:

```

public class MaxStockPriceDriver {

    public static void main(String[] args) throws Exception {

        Configuration configuration = new Configuration();

        Job job = Job.getInstance(configuration, "Calculating max stock price high");

        job.setJarByClass(MaxStockPriceDriver.class);

        job.setMapperClass(MaxStockPriceMap.class);

        job.setCombinerClass(MaxStockPriceReduce.class);

        job.setReducerClass(MaxStockPriceReduce.class);

        job.setOutputKeyClass(Text.class);

        job.setOutputValueClass(DoubleWritable.class);

    }

}

```



```

FileInputFormat.addInputPath(job, new Path(args[0]));

FileOutputFormat.setOutputPath(job, new Path(args[1]));

System.exit(job.waitForCompletion(true) ? 0 : 1);

}

}

```

Run Command:

```
hadoop jar /home/cloudera/Desktop/hw4/mergedMAXPrice.jar mergeMaxStock.MaxStockPriceDriver
/nyseMerged /mergedOutput
```

Run time: 1 minute 8 seconds

```

[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Desktop/hw4/mergedMAXPrice.jar mergeMaxStock.MaxStockPriceDriver /nyseMerged /mergedOutput
22/10/31 06:40:39 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/10/31 06:40:40 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/10/31 06:40:41 INFO input.FileInputFormat: Total input paths to process : 1
22/10/31 06:40:41 INFO mapreduce.JobSubmitter: number of splits:4
22/10/31 06:40:41 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1666102279738_0007
22/10/31 06:40:42 INFO impl.YarnClientImpl: Submitted application application_1666102279738_0007
22/10/31 06:40:42 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1666102279738_0007/
22/10/31 06:40:42 INFO mapreduce.Job: Running job: job_1666102279738_0007
22/10/31 06:40:51 INFO mapreduce.Job: Job job_1666102279738_0007 running in uber mode : false
22/10/31 06:40:51 INFO mapreduce.Job: map 0% reduce 0%
22/10/31 06:41:20 INFO mapreduce.Job: map 7% reduce 0%
22/10/31 06:41:23 INFO mapreduce.Job: map 16% reduce 0%
22/10/31 06:41:24 INFO mapreduce.Job: map 26% reduce 0%
22/10/31 06:41:27 INFO mapreduce.Job: map 35% reduce 0%
22/10/31 06:41:29 INFO mapreduce.Job: map 60% reduce 0%
22/10/31 06:41:30 INFO mapreduce.Job: map 67% reduce 0%
22/10/31 06:41:31 INFO mapreduce.Job: map 75% reduce 0%
22/10/31 06:41:33 INFO mapreduce.Job: map 83% reduce 0%
22/10/31 06:41:34 INFO mapreduce.Job: map 100% reduce 0%
22/10/31 06:41:40 INFO mapreduce.Job: map 100% reduce 100%
22/10/31 06:41:41 INFO mapreduce.Job: Job job_1666102279738_0007 completed successfully
22/10/31 06:41:41 INFO mapreduce.Job: Cleanup: 40

```

Merged Output:

part-r-00000	
part-r-00000(1)	
nyseMerged	
part-r-00000(4)	
AA	94.62
AAI	57.88
AAN	35.21
AAP	83.65
AAR	25.25
AAV	24.78
AB	94.94
ABA	27.94
ABB	33.39
ABC	84.35
ABD	28.58
ABG	30.06
ABK	96.1
ABM	41.63
ABR	34.45
ABT	93.37
ABV	107.5
ABVT	100.0
ABX	54.74
ACC	37.0
ACE	104.0
ACF	64.9
ACG	12.63
ACH	111.6
ACI	112.89
ACL	178.56
ACM	38.25
ACN	44.03
ACO	42.7
ACS	109.55
ACV	65.32
ADC	37.7
ADI	185.5
ADM	48.95
ADP	84.31
ADS	80.79
ADX	40.56
ADY	44.0
AEA	23.94
AEB	26.5
AEC	17.6
AED	26.12
AEE	56.77
AEF	27.0
AEH	148.32
AEH	26.64
AEL	14.6
AEM	83.45

MapReduce on a single file: 1 minute 8 seconds

MapReduce on a bunch of files: 2 minutes 8 seconds

Hence, MapReduce on a single file runs faster than that of a bunch of files.

Part-4: Programming Assignment

Write one MapReduce program using each of the classes that extend FileInputFormat<k,v>

1. CombineFileInputFormat

An abstract InputFormat that returns CombineFileSplit's in InputFormat.getSplits(JobContext) method. Splits are constructed from the files under the input paths. A split cannot have files from different pools. Each split returned may contain blocks from different files. If a maxSplitSize is specified, then blocks on the same node are combined to form a

single split. Blocks that are left over are then combined with other blocks in the same rack. If `maxSplitSize` is not specified, then blocks from the same rack are combined in a single split; no attempt is made to create node-local splits.

Mapper Class:

```
package CombinedFileInputFormat;

import java.io.IOException;
import java.nio.charset.StandardCharsets;
import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class CombinedMap extends Mapper<Object, DoubleWritable, Text, DoubleWritable> {
    IntWritable one = new IntWritable(1);

    public void map(Object key, DoubleWritable value, Context context) throws IOException,
        InterruptedException {
        Text word = new Text();
        context.write(key, one);
    }
}
```

Reducer Class:

```
package CombinedFileInputFormat;

import java.io.IOException;
import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

/**
 *
 * @author zeeni
 */
public class CombinedReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    protected void reduce(Text key, Iterable< DoubleWritable > values,
        Reducer<Text, DoubleWritable, Text, DoubleWritable >.Context context) throws IOException,
        InterruptedException {
        IntWritable finalSum = new IntWritable(0);
        int sum = 0;
        for(DoubleWritable val:values) {
```

```

        sum += val.get();
    }
    finalSum.set(sum);
    System.out.println("Reducing key: " + key + " final sum: " + sum);
    context.write(key, finalSum);
}

}

```

Driver Class:

```

package CombinedFileInputFormat;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.CombineFileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

/**
 *
 * @author zeeni
 */
public class CombinedDriver {
    public static void main(String[] args) throws Exception {
        Configuration configuration = new Configuration();
        //configuration.setInt(FixedLengthInputFormat.FIXED_RECORD_LENGTH, 10);
        configuration.setInt("fixedlengthinputformat.record.length", 8);
        Job job = Job.getInstance(configuration, "Fixed length Input Format");

        job.setJarByClass(CombinedDriver.class);
        job.setInputFormatClass(CombineFileInputFormat.class);

        job.setMapperClass(CombinedMap.class);
        job.setCombinerClass(CombinedReducer.class);
        job.setReducerClass(CombinedReducer.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
    }
}

```

Run Command:


```
hadoop jar /home/cloudera/Desktop/hw4/fixedLength.jar -D
mapreduce.input.fileinputformat.split.maxsize=55 CombinedFileInputFormat.CombinedDriver
/combinedInputOutput
```

Output:

```
22/18/38 22:24:05 INFO ClientRMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/18/38 22:24:06 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/18/38 22:24:06 INFO InputFileInputFormat: Total input paths to process : 1
22/18/38 22:24:06 INFO mapreduce.JobSubmitter: number of splits:5
22/18/38 22:24:06 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1666162279738_0001
22/18/38 22:24:07 INFO ImplVarClienImpl: Submitted application application_1666162279738_0001/
22/18/38 22:24:07 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8080/progress/application_1666162279738_0001/
22/18/38 22:24:07 INFO mapreduce.Job: Running job: job_1666162279738_0001
22/18/38 22:24:15 INFO mapreduce.Job: Job job_1666162279738_0001 running in user mode : false
22/18/38 22:24:28 INFO mapreduce.Job: map 100% reduce 0%
22/18/38 22:24:36 INFO mapreduce.Job: map 100% reduce 100%
22/18/38 22:24:37 INFO mapreduce.Job: Job job_1666162279738_0001 completed successfully
22/18/38 22:24:38 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=789513
    FILE: Number of bytes written=348653
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=3487779
```

2. FixedLengthInputFormat

Input Data: Text file containing zipcodes without any separator and we try to find the count of each zipcode

```
File Edit View 
0212002121021220212302120021210212202123021200212102122021230212002121021220212302120021210212202123021
0212002121021220212302120021210212202123021200212102122021230212002121021220212302120021210212202123021
```

FixedLengthInputFormat is an input format used to read input files which contain fixed-length records. The content of a record need not be text. It can be arbitrary binary data.

Mapper Class:

```
import java.io.IOException;
import java.nio.charset.StandardCharsets;
import org.apache.hadoop.io.BytesWritable;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class FixedLengthInputMap extends Mapper<Object, BytesWritable, Text, IntWritable> {
    IntWritable one = new IntWritable(1);
    public void map(Object key, BytesWritable value, Context context) throws IOException,
    InterruptedException {
        Text word = new Text();
        word.set(new String(value.getBytes(), StandardCharsets.UTF_8));
```

```

        context.write(word, one);
    }

}

```

Reducer Class:

```

package FixedLengthInputFormat;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class FixedLengthInputReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    protected void reduce(Text key, Iterable<IntWritable> values,
        Reducer<Text, IntWritable, Text, IntWritable>.Context context) throws IOException,
        InterruptedException {
        IntWritable finalSum = new IntWritable(0);
        int sum = 0;
        for(IntWritable val:values) {
            sum += val.get();
        }
        finalSum.set(sum);
        System.out.println("Reducing key: "+ key + " final sum: " + sum);
    }
}

```

Driver Class:

```

package FixedLengthInputFormat;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.FixedLengthInputFormat;

public class FixedLengthInputDriver {
    public static void main(String[] args) throws Exception {
        Configuration configuration = new Configuration();
        configuration.setInt("fixedlengthinputformat.record.length", 5);
    }
}

```

```

Job job = Job.getInstance(configuration, "Fixed length Input Format");
job.setJarByClass(FixedLengthInputDriver.class);
job.setInputFormatClass(FixedLengthInputFormat.class);
job.setMapperClass(FixedLengthInputMap.class);
job.setCombinerClass(FixedLengthInputReducer.class);
job.setReducerClass(FixedLengthInputReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}

}

```

Run Command:

```

hadoop jar /home/cloudera/Desktop/hw4/fixedLength.jar
FixedLengthInputFormat.FixedLengthInputDriver /zipcode /fixedInputOutput

```

Output:

```

[cloudera@quickstart Desktop]$ hadoop jar /home/cloudera/Desktop/hw4/fixedLength.jar FixedLengthInputFormat.FixedLengthInputDriver /zipcode /fixedInputOutput
22/10/31 13:23:29 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/10/31 13:23:29 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to
22/10/31 13:23:30 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:952)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:690)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:879)
22/10/31 13:23:30 INFO input.FileInputFormat: Total input paths to process : 1
22/10/31 13:23:30 INFO mapreduce.JobSubmitter: number of splits:1
22/10/31 13:23:30 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1666102279738_0010
22/10/31 13:23:30 INFO impl.YarnClientImpl: Submitted application application_1666102279738_0010
22/10/31 13:23:30 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1666102279738_0010/
22/10/31 13:23:30 INFO mapreduce.Job: Running job: job_1666102279738_0010
22/10/31 13:23:36 INFO mapreduce.Job: Job job_1666102279738_0010 running in uber mode : false
22/10/31 13:23:36 INFO mapreduce.Job:  map 0% reduce 0%

```

/								Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
drwxrwxrwx	hdfs	supergroup	0 B	Wed Jul 19 05:34:46 -0700 2017	0	0 B	benchmarks	
drwxr-xr-x	cloudera	supergroup	0 B	Tue Sep 20 12:58:19 -0700 2022	0	0 B	data	
drwxr-xr-x	cloudera	supergroup	0 B	Mon Oct 31 13:23:53 -0700 2022	0	0 B	fixedInputOutput	
drwxr-xr-x	hbase	supergroup	0 B	Tue Oct 18 09:59:14 -0700 2022	0	0 B	hbase	
drwxr-xr-x	cloudera	supergroup	0 B	Mon Oct 31 00:04:43 -0700 2022	0	0 B	hw4	

3. KeyValueTextInputFormat

An InputFormat for plain text files. Files are broken into lines. Either line feed or carriage-return are used to signal end of line. Each line is divided into key and value parts by a separator byte. If no such a byte exists, the key will be the entire line and value will be empty. The separator byte can be specified in config file under the attribute name

mapreduce.input.keyvaluelinerecordreader.key.value.separator. The default is the tab character ('\t').

Mapper Class:

```
public class KeyValueInputFormatMap extends Mapper<Text, Text, Text,
IntWritable> {
    public void map(Text key, Text value, Context context) throws
IOException, InterruptedException {
        if (key.toString().length() == 1 || value.toString().length() == 1 )
        {
            return;
        } else {
            IntWritable count = new
            IntWritable(Integer.parseInt(value.toString()));
            context.write(key, count);
        }
    }
}
```

Reducer Class:

```
public class KeyValueInputReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    protected void reduce(Text key, Iterable<IntWritable> values,
        Reducer<Text, IntWritable, Text, IntWritable>.Context context) throws
        IOException, InterruptedException {
        IntWritable finalSum = new IntWritable(0);
        int sum = 0;
        for(IntWritable val:values) {
            sum += val.get();
        }
        finalSum.set(sum);
        System.out.println("Reducing key: "+ key + " final sum: " + sum);
        context.write(key, finalSum);
    }
}
```

Driver Class:

```
public class KeyValueInputFormatDriver {
    public static void main(String[] args) throws Exception {
        Configuration configuration = new Configuration();
        configuration.set("mapreduce.input.keyvaluelinerecordreader.key.value.separat
or", " ");
        Job job = Job.getInstance(configuration, "Key value Input Format");
        job.setJarByClass(KeyValueInputFormatDemoMapper.class);
        job.setInputFormatClass(KeyValueTextInputFormat.class);
    }
}
```



```

        job.setMapperClass(KeyValueInputFormatDemoMapper.class);
        job.setCombinerClass(KeyValueInputFormatDemoReducer.class);
        job.setReducerClass(KeyValueInputFormatDemoReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }

}

```

```

20 22:09:48 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=74
  FILE: Number of bytes written=237929
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=698
  HDFS: Number of bytes written=63
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=4469
  Total time spent by all reduces in occupied slots (ms)=4334
  Total time spent by all map tasks (ms)=4469
  Total time spent by all reduce tasks (ms)=4334
  Total vcore-milliseconds taken by all map tasks=4469
  Total vcore-milliseconds taken by all reduce tasks=4334
  Total megabyte-milliseconds taken by all map tasks=4676256
  Total megabyte-milliseconds taken by all reduce tasks=4438016
Map-Reduce Framework
  Map input records=30
  Map output records=24
  Map output bytes=492
  Map output materialized bytes=74
  Input split bytes=106
  Combine input records=24
  Combine output records=3
  Reduce input groups=3
  Reduce shuffle bytes=74
  Reduce input records=3
  Reduce output records=3
  Spilled Records=6
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=104
  CPU time spent (ms)=0
  Physical memory (bytes) snapshot=0
  Virtual memory (bytes) snapshot=0
  Total committed heap usage (bytes)=330825728
Shuffle Errors
  BAD_ID=0

```

4. NLineInputFormat

NLineInputFormat which splits N lines of input as one split.

Mapper Class:

```

public class CombinedMyFormat extends CombineFileInputFormat<LongWritable,Text>{

    @Override
    public RecordReader<LongWritable, Text>
        createRecordReader(InputSplit split, TaskAttemptContext context)
        throws IOException {

```

```

CombineFileRecordReader<LongWritable, Text> reader =
    new CombineFileRecordReader<LongWritable, Text>(
        (CombineFileSplit) split, context, myCombineFileRecordReader.class);
return reader;
}

public static class myCombineFileRecordReader extends RecordReader<LongWritable, Text> {
    private LineRecordReader lineRecordReader = new LineRecordReader();

    public myCombineFileRecordReader(CombineFileSplit split,
        TaskAttemptContext context, Integer index) throws IOException {

        FileSplit fileSplit = new FileSplit(split.getPath(index),
            split.getOffset(index),
            split.getLength(index),
            split.getLocations());
        lineRecordReader.initialize(fileSplit, context);
    }

    @Override
    public void initialize(InputSplit inputSplit, TaskAttemptContext context)
        throws IOException, InterruptedException {
        //linerecordReader.initialize(inputSplit, context);
    }

    @Override
    public void close() throws IOException {
        lineRecordReader.close();
    }

    @Override
    public float getProgress() throws IOException {
        return lineRecordReader.getProgress();
    }

    @Override
    public LongWritable getCurrentKey() throws IOException,
        InterruptedException {
        return lineRecordReader.getCurrentKey();
    }

    @Override
    public Text getCurrentValue() throws IOException, InterruptedException {
        return lineRecordReader.getCurrentValue();
    }
}

```

```

        @Override
        public boolean nextKeyValue() throws IOException, InterruptedException {
            return lineRecordReader.nextKeyValue();
        }
    }
}

```

Driver Class:

```

public class InputOutputDriver {

    public static void main(String args[]) throws IllegalArgumentException, IOException,
    ClassNotFoundException, InterruptedException {
        Configuration conf = new Configuration();
        conf.set("mapreduce.input.keyvaluelinerecordreader.key.value.separator", " ");

        Job job = new Job(conf);
        Job job = Job.getInstance();
        job.getConfiguration().setInt("mapreduce.input.lineinputformat.linespermap", 4);
        job.setJarByClass(HitMain.class);

        // Specify various job-specific parameters
        job.setMapperClass(HitCounter.class);
        job.setReducerClass(HitReducer.class);

        job.setInputFormatClass(NLineInputFormat.class);
    }
}

```

Mapper Class:

```

public class InputOutputMapper extends Mapper<LongWritable,Text,Text,IntWritable>{ //-
NLineInputFormat

    private final static IntWritable one = new IntWritable(1);
    private Text ipaddress = new Text();

    public void map(BytesWritable key, Text value,
org.apache.hadoop.mapreduce.Mapper.Context context) throws IOException, InterruptedException
{ //- NLineInputFormat

        //String line=value.toString();
        String line = new String(value.get(),StandardCharsets.UTF_8);

        String[] tokens=line.split(" ");
        ipaddress.set(tokens[0]);
    }
}

```

```

        context.write(ipaddress, one);
        //context gives output key and value
    }
}

```

5. SequenceFileInputFormat

An InputFormat for SequenceFiles.

6. TextInputFormat

An InputFormat for plain text files. Files are broken into lines. Either linefeed or carriage-return are used to signal end of line. Keys are the position in the file, and values are the line of text.

Mapper Class:

```

public class MaxStockPriceMap extends Mapper<Object, Text, Text, DoubleWritable> {

    protected void map(Object key, Text value, Mapper<Object, Text, Text,
        DoubleWritable>.Context context) throws IOException, InterruptedException {

        String[] tokens = value.toString().split(",");

        Text stockSymbol = new Text();
        stockSymbol.set(tokens[1]);

        DoubleWritable price = new DoubleWritable();

        Double stock_price_high_value = 0.0;

        try {

            stock_price_high_value = Double.parseDouble(tokens[4]);

            price.set(stock_price_high_value);

        }

        catch (NumberFormatException nf) {

            System.out.println("Continue without interruption, its the header line");

        }

        context.write(stockSymbol, price);

    }

}

```

Reducer Class:

```
public class MaxStockPriceReduce extends Reducer<Text, DoubleWritable, Text, DoubleWritable> {  
    public void reduce(Text key, Iterable<DoubleWritable> values, Context context)  
        throws IOException, InterruptedException {  
        double final_stock_high = Double.MIN_VALUE;  
        for (DoubleWritable val: values) {  
            final_stock_high = Math.max(val.get(), final_stock_high);  
        }  
        DoubleWritable final_stock_price_red = new DoubleWritable(final_stock_high);  
        context.write(key, final_stock_price_red);  
    }  
}
```

Driver Class:

```
public class MaxStockPriceDriver {  
    public static void main(String[] args) throws Exception {  
        Configuration configuration = new Configuration();  
        Job job = Job.getInstance(configuration, "Calculating max stock price high");  
        job.setJarByClass(MaxStockPriceDriver.class);  
        job.setMapperClass(MaxStockPriceMap.class);  
        job.setCombinerClass(MaxStockPriceReduce.class);  
        job.setReducerClass(MaxStockPriceReduce.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(DoubleWritable.class);  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

```
}
```

Run Command:

```
hadoop jar /home/cloudera/Desktop/hw4/mergedMAxPrice.jar mergeMaxStock.MaxStockPriceDriver  
/nyseMerged /mergedOutput
```

```
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Desktop/hw4/mergedMAxPrice.jar mergeMaxStock.MaxStockPriceDriver /nyseMerged /mergedOutput
22/10/31 06:40:39 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/10/31 06:40:40 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/10/31 06:40:41 INFO input.FileInputFormat: Total input paths to process : 1
22/10/31 06:40:41 INFO mapreduce.JobSubmitter: number of splits:4
22/10/31 06:40:41 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1666102279738_0007
22/10/31 06:40:42 INFO impl.YarnClientImpl: Submitted application application_1666102279738_0007
22/10/31 06:40:42 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1666102279738_0007/
22/10/31 06:40:42 INFO mapreduce.Job: Running job: job_1666102279738_0007
22/10/31 06:40:51 INFO mapreduce.Job: Job job_1666102279738_0007 running in uber mode : false
22/10/31 06:40:51 INFO mapreduce.Job: map 0% reduce 0%
22/10/31 06:41:20 INFO mapreduce.Job: map 7% reduce 0%
22/10/31 06:41:23 INFO mapreduce.Job: map 16% reduce 0%
22/10/31 06:41:24 INFO mapreduce.Job: map 26% reduce 0%
22/10/31 06:41:27 INFO mapreduce.Job: map 35% reduce 0%
22/10/31 06:41:29 INFO mapreduce.Job: map 60% reduce 0%
22/10/31 06:41:30 INFO mapreduce.Job: map 67% reduce 0%
22/10/31 06:41:31 INFO mapreduce.Job: map 75% reduce 0%
22/10/31 06:41:33 INFO mapreduce.Job: map 83% reduce 0%
22/10/31 06:41:34 INFO mapreduce.Job: map 100% reduce 0%
22/10/31 06:41:40 INFO mapreduce.Job: map 100% reduce 100%
22/10/31 06:41:41 INFO mapreduce.Job: Job job_1666102279738_0007 completed successfully
22/10/31 06:41:41 INFO mapreduce.Job: Cleanup: 40
```

Merged Output:

part-r-00000	part-r-00000(1)	nyseMerged	part-r-00000(4)
AA	94.62		
AAI	57.88		
AAN	35.21		
AAP	83.65		
AAR	25.25		
AAV	24.78		
AB	94.94		
ABA	27.94		
ABB	33.39		
ABC	84.35		
ABD	28.58		
ABG	30.06		
ABK	96.1		
ABM	41.63		
ABR	34.45		
ABT	93.37		
ABV	107.5		
ABVT	100.0		
ABX	54.74		
ACC	37.0		
ACE	104.0		
ACF	64.9		
ACG	12.63		
ACH	111.6		
ACI	112.89		
ACL	178.56		
ACM	38.25		
ACN	44.03		
ACO	42.7		
ACS	109.55		
ACV	65.32		
ADC	37.7		
ADI	185.5		
ADM	48.95		
ADP	84.31		
ADS	80.79		
ADX	40.56		
ADY	44.0		
AEA	23.94		
AEB	26.5		
AEC	17.6		
AED	26.12		
AEE	56.77		
AEF	27.0		
AEH	148.32		
AEI	26.64		
AEL	14.6		
AEM	83.45		

Part-5: Programming Assignment

Create a Writable object that stores some fields from the the NYSE dataset to find

- the date of the max stock_volume
- the date of the min stock_volume
- the max stock_price_adj_close

This will be a custom writable class with the above fields.

Mapper will use this writable object as a value, and Reducer will use this writable object as a value.

Writable Class:

```
public class NyseStock implements Writable {

    private Text maxStockVolDate;

    private Text minStockVolDate;
```

```

private DoubleWritable maxstockVol;

private DoubleWritable minstockVol;

private DoubleWritable stock_price_adj_close;


public NyseStock(Text maxStockVolumeDate, Text
minStockVolumeDate, DoubleWritable maxstockVolume, DoubleWritable
minstockVolume, DoubleWritable stock_price_adj_close) {
    this.maxStockVolDate = maxStockVolumeDate;
    this.minStockVolDate = minStockVolumeDate;
    this.maxstockVol = maxstockVolume;
    this.minstockVol = minstockVolume;
    this.stock_price_adj_close = stock_price_adj_close;
}

public NyseStock() {
    this.maxStockVolDate = new Text();
    this.minStockVolDate = new Text();
    this.maxstockVol = new DoubleWritable();
    this.minstockVol = new DoubleWritable();
    this.stock_price_adj_close = new DoubleWritable();
}

public Text getMaxStockVolDate() {
    return maxStockVolDate;
}

public void setMaxStockVolDate(Text maxStockVolumeDate) {
    this.maxStockVolDate = maxStockVolumeDate;
}

public Text getMinStockVolDate() {
    return minStockVolDate;
}

```



```
}
```

```
public void setMinStockVolDate(Text minStockVolumeDate) {
```

```
    this.minStockVolDate = minStockVolumeDate;
```

```
}
```

```
public DoubleWritable getMaxstockVol() {
```

```
    return maxstockVol;
```

```
}
```

```
public void setMaxstockVol(DoubleWritable stockVolume) {
```

```
    this.maxstockVol = stockVolume;
```

```
}
```

```
public DoubleWritable getMinstockVol() {
```

```
    return minstockVol;
```

```
}
```

```
public void setMinstockVol(DoubleWritable stockVolume) {
```

```
    this.minstockVol = stockVolume;
```

```
}
```

```
public DoubleWritable getStock_price_adj_close() {
```

```
    return stock_price_adj_close;
```

```
}
```

```
public void setStock_price_adj_close(DoubleWritable stock_price_adj_close) {
```

```
    this.stock_price_adj_close = stock_price_adj_close;
```

```
}
```

```

@Override

public void write(DataOutput dataOutput) throws IOException {

    minStockVolDate.write(dataOutput);

    maxStockVolDate.write(dataOutput);

    stock_price_adj_close.write(dataOutput);

    maxstockVol.write(dataOutput);

    minstockVol.write(dataOutput);

}

```

```

@Override

public void readFields(DataInput dataInput) throws IOException {

    minStockVolDate.readFields(dataInput);

    maxStockVolDate.readFields(dataInput);

    stock_price_adj_close.readFields(dataInput);

    maxstockVol.readFields(dataInput);

    minstockVol.readFields(dataInput);

}

```

```

@Override

public String toString() {

return "NyseStockWritable{" +

    "maxStockVolumeDate: " + maxStockVolDate +

    ", minStockVolumeDate: " + minStockVolDate +

    ", maxstockVolume: " + maxstockVol +

    ", minstockVolume: " + minstockVol +

    ", stock_price_adj_close: " + stock_price_adj_close +

    '}';

}

}

```

Mapper Class:

```
public class NyseStockMap extends Mapper<Object, Text, Text, NyseStock> {  
  
    protected void map(Object key, Text value, Mapper<Object, Text, Text, NyseStock>.Context context)  
        throws IOException, InterruptedException  
    {  
        String[] tokens = value.toString().split(",");  
  
        Text stockSymbol = new Text();  
  
        stockSymbol.set(tokens[1]);  
  
        Text maxStockVolumeDate = new Text(tokens[2]);  
  
        try {  
            DoubleWritable stockVolume = new  
                DoubleWritable(Double.parseDouble(tokens[7]));  
  
            DoubleWritable stock_price_adj_close = new  
                DoubleWritable(Double.parseDouble(tokens[8]));  
  
            NyseStock nyseStockWritable = new  
                NyseStock(maxStockVolumeDate, maxStockVolumeDate, stockVolume,  
                    stockVolume, stock_price_adj_close);  
  
            context.write(stockSymbol, nyseStockWritable);  
        }  
  
        catch (NumberFormatException nf) {  
            System.out.println("Continue without interruption, its the headerline");  
        }  
    }  
}
```

Reducer Class:

```
public class NyseStockReduce extends Reducer<Text, NyseStock, Text, NyseStock> {
```

```

public void reduce(Text key, Iterable<NyseStock> values, Context context)
throws IOException, InterruptedException {
    Text maxDate = new Text();
    Text minDate = new Text();
    DoubleWritable maxStock = new DoubleWritable(Double.MIN_VALUE);
    DoubleWritable minStock = new DoubleWritable(Double.MAX_VALUE);
    DoubleWritable stockClosePrice = new
    DoubleWritable(Double.MIN_VALUE);
    for (NyseStock val: values) {
        if (val.getMaxstockVol().get() > maxStock.get()) {
            maxStock.set(val.getMaxstockVol().get());
            maxDate.set(val.getMaxStockVolDate().toString());
        }

        if (val.getMinstockVol().get() < minStock.get()) {
            minStock.set(val.getMinstockVol().get());
            minDate.set(val.getMinStockVolDate().toString());
        }

        if (val.getStock_price_adj_close().get() > stockClosePrice.get())
        {
            stockClosePrice.set(val.getStock_price_adj_close().get());
        }

    }

    NyseStock nyseStockWritable = new NyseStock(maxDate,
    minDate, maxStock, minStock, stockClosePrice);
    context.write(key, nyseStockWritable);
}

```

```
}
```

Driver Class:

```
public class NyseStockDriver {  
  
    public static void main(String[] args) throws IOException, InterruptedException,  
    ClassNotFoundException {  
  
        Configuration configuration = new Configuration();  
  
        Job job = Job.getInstance(configuration, "Custom Writable");  
  
        job.setJarByClass(NyseStockDriver.class);  
  
        job.setMapperClass(NyseStockMap.class);  
  
        job.setReducerClass(NyseStockReduce.class);  
  
  
        job.setOutputKeyClass(Text.class);  
  
        job.setOutputValueClass(NyseStock.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
  
    }  
}
```

Run Command:

```
hadoop jar /home/cloudera/Desktop/hw4/writable.jar NyseWritable.NyseStockDriver /nyseMerged  
/writableOutput
```

Run Time: 1 minute 19 seconds

```
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Desktop/hw4/writable.jar NyseWritable.NyseStockDriver /nyseMerged /writableOutput
22/10/31 09:07:14 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/10/31 09:07:14 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/10/31 09:07:14 INFO Input.FileInputFormat: Total input paths to process : 1
22/10/31 09:07:14 INFO mapreduce.JobSubmitter: number of splits:4
22/10/31 09:07:15 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1666102279738_0008
22/10/31 09:07:15 INFO impl.YarnClientImpl: Submitted application application_1666102279738_0008
22/10/31 09:07:15 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1666102279738_0008/
22/10/31 09:07:15 INFO mapreduce.Job: Running job: job_1666102279738_0008
22/10/31 09:07:21 INFO mapreduce.Job: Job job_1666102279738_0008 running in uber mode : false
22/10/31 09:07:21 INFO mapreduce.Job: map 0% reduce 0%
22/10/31 09:07:44 INFO mapreduce.Job: map 6% reduce 0%
22/10/31 09:07:49 INFO mapreduce.Job: map 13% reduce 0%
22/10/31 09:07:51 INFO mapreduce.Job: map 18% reduce 0%
22/10/31 09:07:55 INFO mapreduce.Job: map 36% reduce 0%
22/10/31 09:07:56 INFO mapreduce.Job: map 40% reduce 0%
22/10/31 09:07:57 INFO mapreduce.Job: map 44% reduce 0%
22/10/31 09:08:02 INFO mapreduce.Job: map 51% reduce 0%
22/10/31 09:08:03 INFO mapreduce.Job: map 58% reduce 0%
22/10/31 09:08:08 INFO mapreduce.Job: map 78% reduce 0%
22/10/31 09:08:12 INFO mapreduce.Job: map 80% reduce 0%
22/10/31 09:08:13 INFO mapreduce.Job: map 92% reduce 0%
22/10/31 09:08:14 INFO mapreduce.Job: map 100% reduce 0%
22/10/31 09:08:30 INFO mapreduce.Job: map 100% reduce 100%
22/10/31 09:08:31 INFO mapreduce.Job: Job job_1666102279738_0008 completed successfully
22/10/31 09:08:32 INFO mapreduce.Job: Counters: 49
```

Driver Name	Driver Type	Configuration	Start Time	End Time	Status	User
drwxr-xr-x	hdfs	supergroup	0 B	Wed Oct 19 07:23:03 -0700 2022	0	0 B user
drwxr-xr-x	hdfs	supergroup	0 B	Wed Jul 19 05:36:28 -0700 2017	0	0 B var
drwxr-xr-x	cloudera	supergroup	0 B	Mon Oct 31 09:08:29 -0700 2022	0	0 B writableOutput

Output:

part-r-00000	part-r-00000(1)	nyseMerged	part-r-00000(4)	part-r-00000(5)
AA	NyseStockWritable(maxStockVolumeDate: 2009-03-19, minStockVolumeDate: 1965-06-21, maxstockVolume: 2.421065E8, minstockVolume: 0.0, stock price adj close: 44.18)			
AAI	NyseStockWritable(maxStockVolumeDate: 2009-10-08, minStockVolumeDate: 1994-09-14, maxstockVolume: 3.0579E7, minstockVolume: 9600.0, stock price adj close: 33.5)			
AAN	NyseStockWritable(maxStockVolumeDate: 2009-06-08, minStockVolumeDate: 1996-07-05, maxstockVolume: 6602800.0, minstockVolume: 200.0, stock price adj close: 34.36)			
AAP	NyseStockWritable(maxStockVolumeDate: 2006-06-29, minStockVolumeDate: 2002-02-19, maxstockVolume: 1.40877E7, minstockVolume: 9900.0, stock price adj close: 46.92)			
AAR	NyseStockWritable(maxStockVolumeDate: 2003-04-15, minStockVolumeDate: 2000-07-27, maxstockVolume: 475200.0, minstockVolume: 500.0, stock price adj close: 21.5)			
AAV	NyseStockWritable(maxStockVolumeDate: 2006-01-24, minStockVolumeDate: 2004-08-02, maxstockVolume: 5011000.0, minstockVolume: 0.0, stock price adj close: 13.3)			
AB	NyseStockWritable(maxStockVolumeDate: 2009-04-23, minStockVolumeDate: 1909-05-26, maxstockVolume: 3258200.0, minstockVolume: 7600.0, stock price adj close: 79.18)			
ABA	NyseStockWritable(maxStockVolumeDate: 2006-07-10, minStockVolumeDate: 2007-07-05, maxstockVolume: 355200.0, minstockVolume: 100.0, stock price adj close: 27.0)			
ABB	NyseStockWritable(maxStockVolumeDate: 2008-02-13, minStockVolumeDate: 2002-07-09, maxstockVolume: 2.86948E7, minstockVolume: 400.0, stock price adj close: 31.56)			
ABC	NyseStockWritable(maxStockVolumeDate: 2003-09-19, minStockVolumeDate: 1996-02-12, maxstockVolume: 5.5356E7, minstockVolume: 3200.0, stock price adj close: 28.55)			
ABD	NyseStockWritable(maxStockVolumeDate: 2008-11-18, minStockVolumeDate: 2009-12-24, maxstockVolume: 5868000.0, minstockVolume: 43200.0, stock price adj close: 28.22)			
ABG	NyseStockWritable(maxStockVolumeDate: 2006-09-19, minStockVolumeDate: 2002-12-24, maxstockVolume: 4162900.0, minstockVolume: 3900.0, stock price adj close: 27.64)			
ABK	NyseStockWritable(maxStockVolumeDate: 2009-08-27, minStockVolumeDate: 1996-01-08, maxstockVolume: 1.128342E8, minstockVolume: 6600.0, stock price adj close: 93.59)			
ABM	NyseStockWritable(maxStockVolumeDate: 1988-06-08, minStockVolumeDate: 1993-07-29, maxstockVolume: 2388000.0, minstockVolume: 0.0, stock price adj close: 28.41)			
ABR	NyseStockWritable(maxStockVolumeDate: 2009-06-26, minStockVolumeDate: 2004-09-23, maxstockVolume: 2115600.0, minstockVolume: 1500.0, stock price adj close: 25.87)			
ABT	NyseStockWritable(maxStockVolumeDate: 1988-06-03, minStockVolumeDate: 1994-02-28, maxstockVolume: 2.75488E7, minstockVolume: 158200.0, stock price adj close: 57.02)			
ABV	NyseStockWritable(maxStockVolumeDate: 2004-03-03, minStockVolumeDate: 1997-02-28, maxstockVolume: 8896800.0, minstockVolume: 100.0, stock price adj close: 106.09)			
ABVT	NyseStockWritable(maxStockVolumeDate: 2009-06-26, minStockVolumeDate: 2004-08-19, maxstockVolume: 1538000.0, minstockVolume: 0.0, stock price adj close: 66.51)			
ABX	NyseStockWritable(maxStockVolumeDate: 2009-09-09, minStockVolumeDate: 1985-05-30, maxstockVolume: 5.3779E7, minstockVolume: 2400.0, stock price adj close: 52.4)			
ACC	NyseStockWritable(maxStockVolumeDate: 2009-05-06, minStockVolumeDate: 2005-05-27, maxstockVolume: 4164400.0, minstockVolume: 1900.0, stock price adj close: 31.47)			
ACE	NyseStockWritable(maxStockVolumeDate: 2008-07-17, minStockVolumeDate: 1995-02-22, maxstockVolume: 7.44371E7, minstockVolume: 0.0, stock price adj close: 63.68)			
ACF	NyseStockWritable(maxStockVolumeDate: 2003-01-16, minStockVolumeDate: 1992-06-12, maxstockVolume: 4.42225E7, minstockVolume: 6200.0, stock price adj close: 63.63)			
ACG	NyseStockWritable(maxStockVolumeDate: 1987-08-21, minStockVolumeDate: 2009-11-04, maxstockVolume: 7523000.0, minstockVolume: 5000.0, stock price adj close: 8.25)			
ACH	NyseStockWritable(maxStockVolumeDate: 2007-09-12, minStockVolumeDate: 2006-12-06, maxstockVolume: 1.15053E7, minstockVolume: 0.0, stock price adj close: 86.77)			
ACI	NyseStockWritable(maxStockVolumeDate: 2009-07-28, minStockVolumeDate: 1994-10-26, maxstockVolume: 2.4998E7, minstockVolume: 200.0, stock price adj close: 73.29)			
ACL	NyseStockWritable(maxStockVolumeDate: 2010-01-04, minStockVolumeDate: 2003-11-28, maxstockVolume: 1.38223E7, minstockVolume: 72300.0, stock price adj close: 109.14)			
ACM	NyseStockWritable(maxStockVolumeDate: 2007-05-10, minStockVolumeDate: 2007-11-23, maxstockVolume: 2.08079E7, minstockVolume: 158800.0, stock price adj close: 37.25)			
ACN	NyseStockWritable(maxStockVolumeDate: 2009-08-31, minStockVolumeDate: 2001-12-24, maxstockVolume: 6.74614E7, minstockVolume: 181100.0, stock price adj close: 43.75)			
ACO	NyseStockWritable(maxStockVolumeDate: 2006-07-21, minStockVolumeDate: 1992-01-21, maxstockVolume: 2555200.0, minstockVolume: 400.0, stock price adj close: 38.49)			
ACS	NyseStockWritable(maxStockVolumeDate: 2001-06-26, minStockVolumeDate: 2009-11-25, maxstockVolume: 3.748E7, minstockVolume: 0.0, stock price adj close: 63.92)			
ACV	NyseStockWritable(maxStockVolumeDate: 2006-11-16, minStockVolumeDate: 1984-12-24, maxstockVolume: 9806700.0, minstockVolume: 0.0, stock price adj close: 29.92)			
ADC	NyseStockWritable(maxStockVolumeDate: 2003-07-30, minStockVolumeDate: 2000-12-12, maxstockVolume: 659600.0, minstockVolume: 0.0, stock price adj close: 27.87)			
ADI	NyseStockWritable(maxStockVolumeDate: 2006-08-11, minStockVolumeDate: 1990-03-13, maxstockVolume: 2.7748E7, minstockVolume: 13800.0, stock price adj close: 89.87)			
ADM	NyseStockWritable(maxStockVolumeDate: 2007-02-16, minStockVolumeDate: 1984-06-11, maxstockVolume: 5.24148E7, minstockVolume: 61700.0, stock price adj close: 46.6)			
ADP	NyseStockWritable(maxStockVolumeDate: 2003-03-13, minStockVolumeDate: 2002-07-09, maxstockVolume: 2.48474E7, minstockVolume: 99200.0, stock price adj close: 52.67)			
ADS	NyseStockWritable(maxStockVolumeDate: 2008-01-28, minStockVolumeDate: 2001-08-06, maxstockVolume: 3.40658E7, minstockVolume: 2700.0, stock price adj close: 80.72)			
ADX	NyseStockWritable(maxStockVolumeDate: 2002-01-02, minStockVolumeDate: 1984-10-17, maxstockVolume: 1910800.0, minstockVolume: 400.0, stock price adj close: 13.48)			
ADY	NyseStockWritable(maxStockVolumeDate: 2009-07-13, minStockVolumeDate: 2004-06-02, maxstockVolume: 3238400.0, minstockVolume: 0.0, stock price adj close: 43.16)			
AEA	NyseStockWritable(maxStockVolumeDate: 2005-03-02, minStockVolumeDate: 2005-01-07, maxstockVolume: 4328000.0, minstockVolume: 29200.0, stock price adj close: 17.69)			
AEB	NyseStockWritable(maxStockVolumeDate: 2006-09-06, minStockVolumeDate: 2006-07-08, maxstockVolume: 1277100.0, minstockVolume: 2000.0, stock price adj close: 20.07)			
AEC	NyseStockWritable(maxStockVolumeDate: 1996-10-15, minStockVolumeDate: 1994-11-25, maxstockVolume: 2417500.0, minstockVolume: 200.0, stock price adj close: 13.67)			
AED	NyseStockWritable(maxStockVolumeDate: 2005-11-30, minStockVolumeDate: 2005-11-25, maxstockVolume: 1872200.0, minstockVolume: 0.0, stock price adj close: 19.59)			
AEE	NyseStockWritable(maxStockVolumeDate: 2009-09-10, minStockVolumeDate: 1998-12-24, maxstockVolume: 1.76048E7, minstockVolume: 31000.0, stock price adj close: 47.91)			
AEF	NyseStockWritable(maxStockVolumeDate: 2007-09-20, minStockVolumeDate: 2007-09-19, maxstockVolume: 1268200.0, minstockVolume: 0.0, stock price adj close: 21.52)			
AEG	NyseStockWritable(maxStockVolumeDate: 2008-09-18, minStockVolumeDate: 1986-11-11, maxstockVolume: 6098300.0, minstockVolume: 0.0, stock price adj close: 44.33)			
AEN	NyseStockWritable(maxStockVolumeDate: 2005-06-15, minStockVolumeDate: 2006-07-03, maxstockVolume: 2748700.0, minstockVolume: 7000.0, stock price adj close: 19.62)			
AEL	NyseStockWritable(maxStockVolumeDate: 2004-12-01, minStockVolumeDate: 2004-11-26, maxstockVolume: 5985900.0, minstockVolume: 9600.0, stock price adj close: 13.96)			
AEM	NyseStockWritable(maxStockVolumeDate: 2009-10-30, minStockVolumeDate: 1985-09-27, maxstockVolume: 1.5153E7, minstockVolume: 300.0, stock price adj close: 79.74)			

Part-6: Programming Assignment

Redo Part5 of this assignment, but cram multiple values (max stock_volume, min stock_volume, max stock_price_adj_close) into a Text object with some delimiter. Use a Combiner. Compare the running time of Part 2 to Part 3. You could measure the running time programmatically, or use your smartphone's timer.

Driver Class with Combiner:

```
public class NyseStockDriverCombiner {
```

```

public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException
{
    Configuration configuration = new Configuration();

    Job job = Job.getInstance(configuration, "Custom Writable");

    job.setJarByClass(NyseStockDriver.class);

    job.setMapperClass(NyseWritable.NyseStockMap.class);

    job.setCombinerClass(NyseWritable.NyseStockReduce.class);

    job.setReducerClass(NyseWritable.NyseStockReduce.class);


    job.setOutputKeyClass(Text.class);

    job.setOutputValueClass(NyseWritable.NyseStock.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));

    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);

}
}

```

Run Command:

```

hadoop jar /home/cloudera/Desktop/hw4/writableCombiner.jar
NyseWritableCombiner.NyseStockDriverCombiner /nyseMerged /writableCombinerOutput

```

Run Time: 50 seconds

```

[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Desktop/hw4/writableCombiner.jar NyseWritableCombiner.NyseStockDriverCombiner /nyseMerged /writableCombinerOutput
22/10/31 10:02:33 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/10/31 10:02:33 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/10/31 10:02:34 INFO input.FileInputFormat: Total input paths to process : 1
22/10/31 10:02:34 INFO mapreduce.JobSubmitter: number of splits:4
22/10/31 10:02:34 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1666102279738_0009
22/10/31 10:02:34 INFO impl.YarnClientImpl: Submitted application application_1666102279738_0009
22/10/31 10:02:34 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1666102279738_0009/
22/10/31 10:02:34 INFO mapreduce.Job: Running job: job_1666102279738_0009
22/10/31 10:02:40 INFO mapreduce.Job: Job job_1666102279738_0009 running in uber mode : false
22/10/31 10:02:40 INFO mapreduce.Job: map 0% reduce 0%
22/10/31 10:03:02 INFO mapreduce.Job: map 9% reduce 0%
22/10/31 10:03:06 INFO mapreduce.Job: map 19% reduce 0%
22/10/31 10:03:08 INFO mapreduce.Job: map 33% reduce 0%
22/10/31 10:03:09 INFO mapreduce.Job: map 44% reduce 0%
22/10/31 10:03:12 INFO mapreduce.Job: map 49% reduce 0%
22/10/31 10:03:13 INFO mapreduce.Job: map 60% reduce 0%
22/10/31 10:03:14 INFO mapreduce.Job: map 60% reduce 0%
22/10/31 10:03:15 INFO mapreduce.Job: map 90% reduce 0%
22/10/31 10:03:17 INFO mapreduce.Job: map 100% reduce 0%
22/10/31 10:03:21 INFO mapreduce.Job: map 100% reduce 100%
22/10/31 10:03:21 INFO mapreduce.Job: Job job_1666102279738_0009 completed successfully
22/10/31 10:03:21 INFO mapreduce.Job: Counters: 50
    File System Counters
      FILE: Number of bytes read=297062
      FILE: Number of bytes written=1070600

```

drwxr-xr-x	hdfs	supergroup	0 B	Wed Jul 19 05:36:28 -0700 2017	0	0 B	var
drwxr-xr-x	cloudera	supergroup	0 B	Mon Oct 31 10:03:19 -0700 2022	0	0 B	writableCombinerOutput
drwxr-xr-x	cloudera	supergroup	0 B	Mon Oct 31 09:08:29 -0700 2022	0	0 B	writableOutput

Output:

part-r-00000	part-r-00000(1)	nyseMerged	part-r-00000(4)	part-r-00000(5)	part-r-00000(6)
AA	NyseStockWritable{maxStockVolumeDate: 2009-03-19, minStockVolumeDate: 1965-06-21, maxstockVolume: 2.421065E8, minstockVolume: 0.0, stock price adj close: 44.18}				
AAI	NyseStockWritable{maxStockVolumeDate: 2009-10-08, minStockVolumeDate: 1994-09-14, maxstockVolume: 3.0579E7, minstockVolume: 9600.0, stock price adj close: 33.5}				
AAN	NyseStockWritable{maxStockVolumeDate: 2009-06-08, minStockVolumeDate: 1996-07-05, maxstockVolume: 6602800.0, minstockVolume: 200.0, stock price adj close: 34.36}				
AAP	NyseStockWritable{maxStockVolumeDate: 2006-06-29, minStockVolumeDate: 2002-02-19, maxstockVolume: 1.40077E7, minstockVolume: 9900.0, stock price adj close: 46.92}				
AAR	NyseStockWritable{maxStockVolumeDate: 2003-04-15, minStockVolumeDate: 2000-07-27, maxstockVolume: 475200.0, minstockVolume: 500.0, stock price adj close: 21.5}				
AAV	NyseStockWritable{maxStockVolumeDate: 2006-01-24, minStockVolumeDate: 2004-08-02, maxstockVolume: 5011000.0, minstockVolume: 0.0, stock price adj close: 13.3}				
AB	NyseStockWritable{maxStockVolumeDate: 2009-04-23, minStockVolumeDate: 1989-05-26, maxstockVolume: 3258200.0, minstockVolume: 7600.0, stock price adj close: 79.18}				
ABA	NyseStockWritable{maxStockVolumeDate: 2006-07-10, minStockVolumeDate: 2007-07-05, maxstockVolume: 355200.0, minstockVolume: 100.0, stock price adj close: 27.0}				
ABB	NyseStockWritable{maxStockVolumeDate: 2008-02-13, minStockVolumeDate: 2002-07-09, maxstockVolume: 2.86948E7, minstockVolume: 400.0, stock price adj close: 31.56}				
ABC	NyseStockWritable{maxStockVolumeDate: 2003-09-19, minStockVolumeDate: 1996-02-12, maxstockVolume: 5.5356E7, minstockVolume: 3200.0, stock price adj close: 28.55}				
ABD	NyseStockWritable{maxStockVolumeDate: 2008-11-18, minStockVolumeDate: 2009-12-24, maxstockVolume: 5860000.0, minstockVolume: 43200.0, stock price adj close: 28.22}				
ABG	NyseStockWritable{maxStockVolumeDate: 2006-09-19, minStockVolumeDate: 2002-12-24, maxstockVolume: 4162900.0, minstockVolume: 3900.0, stock price adj close: 27.64}				
ABK	NyseStockWritable{maxStockVolumeDate: 2009-08-27, minStockVolumeDate: 1996-01-08, maxstockVolume: 1.128342E8, minstockVolume: 6600.0, stock price adj close: 93.59}				
ABM	NyseStockWritable{maxStockVolumeDate: 1988-06-08, minStockVolumeDate: 1993-07-29, maxstockVolume: 2388000.0, minstockVolume: 0.0, stock price adj close: 28.41}				
ABR	NyseStockWritable{maxStockVolumeDate: 2009-06-26, minStockVolumeDate: 2004-09-23, maxstockVolume: 2115600.0, minstockVolume: 1500.0, stock price adj close: 25.87}				
ABT	NyseStockWritable{maxStockVolumeDate: 1988-06-03, minStockVolumeDate: 1994-02-28, maxstockVolume: 2.75488E7, minstockVolume: 158200.0, stock price adj close: 57.02}				
ABV	NyseStockWritable{maxStockVolumeDate: 2004-03-03, minStockVolumeDate: 1997-02-28, maxstockVolume: 8896000.0, minstockVolume: 100.0, stock price adj close: 106.09}				
ABVT	NyseStockWritable{maxStockVolumeDate: 2009-06-26, minStockVolumeDate: 2004-08-19, maxstockVolume: 1538000.0, minstockVolume: 0.0, stock price adj close: 66.51}				
ABX	NyseStockWritable{maxStockVolumeDate: 2009-09-09, minStockVolumeDate: 1985-05-30, maxstockVolume: 5.3779E7, minstockVolume: 2400.0, stock price adj close: 52.4}				
ACC	NyseStockWritable{maxStockVolumeDate: 2009-05-06, minStockVolumeDate: 2005-05-27, maxstockVolume: 4164400.0, minstockVolume: 1900.0, stock price adj close: 31.47}				
ACE	NyseStockWritable{maxStockVolumeDate: 2008-07-17, minStockVolumeDate: 1995-02-22, maxstockVolume: 7.44371E7, minstockVolume: 0.0, stock price adj close: 63.68}				
ACF	NyseStockWritable{maxStockVolumeDate: 2003-01-16, minStockVolumeDate: 1992-06-12, maxstockVolume: 4.42225E7, minstockVolume: 6200.0, stock price adj close: 63.63}				
ACG	NyseStockWritable{maxStockVolumeDate: 1987-08-21, minStockVolumeDate: 2009-11-04, maxstockVolume: 7523000.0, minstockVolume: 5000.0, stock price adj close: 8.25}				
ACH	NyseStockWritable{maxStockVolumeDate: 2007-09-12, minStockVolumeDate: 2006-12-06, maxstockVolume: 1.15053E7, minstockVolume: 0.0, stock price adj close: 86.77}				
ACI	NyseStockWritable{maxStockVolumeDate: 2009-07-28, minStockVolumeDate: 1994-10-26, maxstockVolume: 2.4998E7, minstockVolume: 200.0, stock price adj close: 73.29}				
ACL	NyseStockWritable{maxStockVolumeDate: 2010-01-04, minStockVolumeDate: 2003-11-28, maxstockVolume: 1.38223E7, minstockVolume: 72300.0, stock price adj close: 169.14}				
ACM	NyseStockWritable{maxStockVolumeDate: 2007-05-10, minStockVolumeDate: 2007-11-23, maxstockVolume: 2.00079E7, minstockVolume: 158800.0, stock price adj close: 37.25}				
ACN	NyseStockWritable{maxStockVolumeDate: 2009-08-31, minStockVolumeDate: 2001-12-24, maxstockVolume: 6.74614E7, minstockVolume: 181100.0, stock price adj close: 43.75}				
ACO	NyseStockWritable{maxStockVolumeDate: 2006-07-21, minStockVolumeDate: 1992-01-21, maxstockVolume: 2555200.0, minstockVolume: 400.0, stock price adj close: 38.49}				
ACS	NyseStockWritable{maxStockVolumeDate: 2001-06-26, minStockVolumeDate: 1994-11-25, maxstockVolume: 3.748E7, minstockVolume: 0.0, stock price adj close: 63.92}				
ACV	NyseStockWritable{maxStockVolumeDate: 2006-11-16, minStockVolumeDate: 1984-12-24, maxstockVolume: 9806700.0, minstockVolume: 0.0, stock price adj close: 29.92}				
ADC	NyseStockWritable{maxStockVolumeDate: 2003-07-30, minStockVolumeDate: 2000-12-12, maxstockVolume: 659600.0, minstockVolume: 0.0, stock price adj close: 27.87}				
ADI	NyseStockWritable{maxStockVolumeDate: 2006-08-11, minStockVolumeDate: 1990-03-13, maxstockVolume: 2.7748E7, minstockVolume: 13800.0, stock price adj close: 89.07}				
ADM	NyseStockWritable{maxStockVolumeDate: 2007-02-16, minStockVolumeDate: 1984-06-11, maxstockVolume: 5.24148E7, minstockVolume: 61700.0, stock price adj close: 46.6}				
ADP	NyseStockWritable{maxStockVolumeDate: 2003-03-13, minStockVolumeDate: 1983-10-10, maxstockVolume: 2.48474E7, minstockVolume: 99200.0, stock price adj close: 52.67}				
ADS	NyseStockWritable{maxStockVolumeDate: 2008-01-28, minStockVolumeDate: 2001-08-06, maxstockVolume: 3.40658E7, minstockVolume: 2700.0, stock price adj close: 80.72}				
ADX	NyseStockWritable{maxStockVolumeDate: 2002-01-02, minStockVolumeDate: 1984-10-17, maxstockVolume: 1910800.0, minstockVolume: 400.0, stock price adj close: 13.48}				
ADY	NyseStockWritable{maxStockVolumeDate: 2009-07-13, minStockVolumeDate: 2004-06-02, maxstockVolume: 3230400.0, minstockVolume: 0.0, stock price adj close: 43.16}				
AEA	NyseStockWritable{maxStockVolumeDate: 2005-03-02, minStockVolumeDate: 2005-01-07, maxstockVolume: 4328000.0, minstockVolume: 29200.0, stock price adj close: 17.69}				
AEB	NyseStockWritable{maxStockVolumeDate: 2006-09-06, minStockVolumeDate: 2006-07-03, maxstockVolume: 1277100.0, minstockVolume: 2000.0, stock price adj close: 20.07}				
AEC	NyseStockWritable{maxStockVolumeDate: 1996-10-15, minStockVolumeDate: 1994-11-25, maxstockVolume: 2417500.0, minstockVolume: 200.0, stock price adj close: 13.67}				
AED	NyseStockWritable{maxStockVolumeDate: 2005-11-30, minStockVolumeDate: 2005-11-25, maxstockVolume: 1872200.0, minstockVolume: 0.0, stock price adj close: 19.59}				
AEE	NyseStockWritable{maxStockVolumeDate: 2009-09-10, minStockVolumeDate: 1998-12-24, maxstockVolume: 1.76048E7, minstockVolume: 31000.0, stock price adj close: 47.91}				
AEF	NyseStockWritable{maxStockVolumeDate: 2007-09-20, minStockVolumeDate: 2007-09-19, maxstockVolume: 1268200.0, minstockVolume: 0.0, stock price adj close: 21.52}				
AEG	NyseStockWritable{maxStockVolumeDate: 2008-09-18, minStockVolumeDate: 1986-11-11, maxstockVolume: 6698300.0, minstockVolume: 0.0, stock price adj close: 44.33}				
AEH	NyseStockWritable{maxStockVolumeDate: 2005-06-15, minStockVolumeDate: 2006-07-03, maxstockVolume: 2748700.0, minstockVolume: 7000.0, stock price adj close: 19.62}				
AEL	NyseStockWritable{maxStockVolumeDate: 2004-12-01, minStockVolumeDate: 2004-11-26, maxstockVolume: 5985900.0, minstockVolume: 9600.0, stock price adj close: 13.96}				
AEM	NyseStockWritable{maxStockVolumeDate: 2009-10-30, minStockVolumeDate: 1985-09-27, maxstockVolume: 1.5153E7, minstockVolume: 300.0, stock price adj close: 79.74}				

The run time for running the MapReduce without a combiner is 1 minute 19 seconds and that of with combiner is 50 seconds.